



Data Preparation for Machine Learning



Databricks Academy



© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Learning goals

Upon completion of this content, you should be able to:

- Describe Databricks Data Intelligence Platform and its features for machine learning.
- Explain data storage and governance features on Databricks.
- Perform exploratory data analysis and feature engineering using Spark and integrated visualization tools.
- Perform data pre-processing for missing data handling, data encoding and data standardization.
- Use Feature Store for storing and retrieving features.



Prerequisites/Technical Considerations

Things to keep in mind before you work through this course

Prerequisites

1

Familiarity with Databricks workspace and notebooks

2

Familiarity with Delta Lake and Lakehouse

3

Intermediate level knowledge of Python

Technical Considerations

1

Cluster running on **DBR ML 17.3**

2

Unity Catalog enabled workspace



AGENDA

1. Managing and Exploring Data	DEMO	LAB
Managing and Exploring Data in the Lakehouse	✓	✓
02. Data Preparation and Feature Engineering		
Fundamentals of Data Preparation and Feature Engineering		
Data Imputation	✓	✓
Data Encoding		
Data Standardization		
03. Feature Store		
Introduction to Feature Store	✓	✓





Managing and Exploring Data



Data Preparation for Machine Learning



© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Learning objectives

Things you'll be able to do after completing this module

- Understand Databricks Machine Learning as a data-centric platform.
- Explain the benefits of Delta Lake for machine learning, including version control and data integrity.
- Describe the benefits of Unity Catalog for machine learning, including data discovery and collaboration.
- Load and write data from/to Delta tables.
- Use data profiling and visualization to explore and analyze machine learning data.





Managing and Exploring Data

LECTURE

Managing and Exploring Data in the Lakehouse



Lakeflow
Ingest, ETL,
Streaming



Databricks
SQL
Data
Warehousing



AI/BI
Business
Intelligence



AI/ML
Data Science
GenAI



Knowledge of your data



Unity Catalog

Your data stored in an open, broadly
accessible lakehouse format




DELTA LAKE

ICEBERG



Parquet



End-to-end AI


- MLOps (MLflow)
- Model Training
- Model Serving
- Monitoring
- Governance

Gen AI

- Custom models
- Model serving
- RAG

Databricks for Machine Learning

A built-in and collaborative solution for the full ML lifecycle



Features of Databricks for Machine Learning

Non exhaustive list of features that will be used throughout this module

- Collaborative notebooks
- ML Runtime
- Governance of Data & Models (*via Unity Catalog*)
- Feature Store
- Managed MLflow
- Model Training
- Model Serving



Collaborative Multi-Language Notebooks

Collaborative, reproducible, and enterprise ready

Multi-Language


- Use Python, SQL, Scala, and R in one notebook
- Supports: %sh, %fs, %md

Visualizations

- Built-in visualizations
- Support for popular visualization libraries (e.g. *matplotlib*, *ggplot*)

Debug and Optimize

- AI-assisted coding and debugging (Genie Code)
- Get suggestions and explanations



Reproducible

- Track version history automatically
- Use Git Folders for version control

Collaborative

- Share notebooks and add comments
- Collaborate in real time


Interactive

- Use widgets for inputs
- Interact with outputs: Filter data, build dashboards, etc.




Quick Exploratory Data Analysis

Built-in tools for visualizing and understanding data in ML workflow



Create **interactive charts** to visualize data in the Notebook with only two clicks




Summarize a data set's essential properties and statistics in a **data profile** with the push of a button



Tools for Quick ML Model Development

Multi-language support, use standard libraries and custom modules



The screenshot shows a Databricks notebook cell with SQL code. A dropdown menu is open, allowing the user to change the language of the cell. The code in the cell is:

```
1 %sql
2 SELECT * FROM clemens.bikeshare_features LIMIT 10
```


The dropdown menu options are:

- Markdown
- Python (Notebook default)
- Scala
- SQL
- R

Below the code, a table of results is displayed:

	dteday	average_trips
1	2011-01-29	47.73913043478261
2	2012-03-04	142.625
3	2012-10-21	284.3333333333333
4	2011-01-29	47.860565217391305

Mix and match languages based on use case and preferred workflow, choosing from **Python, SQL, Scala, and R**



The screenshot shows two code cells in a Databricks notebook. The first cell contains:

```
1 %pip install folium seaborn==0.11.1
```


The second cell contains:

```
1 import utils
2
3 df2 = utils.scrub(df1, drop="num_columns")
```

Install Python libraries for a notebook without affecting other users with `%pip`

Import local modules using **arbitrary file support** when working in repositories






Unity Catalog


- Context-aware search
- Auto describe tables and columns
- Automated lineage
- End-to-end observability and monitoring
- Sharing AI models

Today, data and AI governance is complex

Data Consumers



Data Governance Team



Unity Catalog (UC)


Unified governance for AI; data, code, models

Unified visibility into data and AI

Single permission model for data and AI


AI-powered monitoring and observability


Open data sharing



The 3-level Namespace of UC

How to use UC






Delta Lake

- Open-source
- Unified data management layer
- Reliable and fast
- Optimization features for storing data in the cloud

What is Delta Lake?

- A **unified data management layer** that brings data reliability and fast analytics to cloud data lakes. It is the optimized storage layer that is the foundation for storing data and tables in the Databricks Data Intelligence Platform.



Delta Lake and Its Features

Open-source, default storage format on Databricks

- Delta Lake is an **open-source** project.
- It is the **default format** for the tables created in Databricks.
- Delta Lake **optimizes performance** with large datasets, with **ACID transactions** and **scalable metadata handling**.
- Designed to improve data reliability, quality, and performance in data lakes.



Delta Lake features

Key features

- Unified **batch** and **streaming**
- Automatic **schema validation**
- Support upserts using the merge operation
- Update your table schema without rewriting data.
- Track row-level changes with Change Data Feed
- **Time-travel**; querying previous versions of a table based on version number or timestamp
- **Performance optimization** with data skipping and liquid clustering
- Supports multiple programming languages like Python, Scala, and SQL.






Managing and Exploring Data

DEMONSTRATION

Load and Explore Data



Demo

Outline

What we'll cover:

- Load data into a Spark DataFrame
 - Read a dataset from a shared volume using Spark
- Explore and understand dataset structure
 - Inspect schema and preview data
 - Compute summary statistics using built-in Spark functions
- Use Genie Code to perform ML-focused exploratory analysis
 - Generate visualizations for target distribution, missing values, and feature relationships
 - Interpret patterns that are useful for machine learning workflows
- Write data to a Delta table
 - Persist a DataFrame as a managed Delta table in Unity Catalog
- Use Delta Lake time travel
 - Read a previous version of a Delta table
 - Understand how versioning supports reproducibility in data workflows





Managing and Exploring Data

LAB EXERCISE

Load and Explore Data



Lab

Outline

What you'll do:

Load data from a CSV file and store it as a Delta table

- Manage access permissions on a Delta table
- Compute and interpret summary statistics
- Use Genie Code to explore the dataset, including:
 - Assessing class balance and target distribution
 - Identifying missing value patterns
 - Exploring feature distributions
 - Comparing key features against the target variable
- Modify the dataset by removing a column
- Use Delta Lake time travel to access a previous version of the data
- View and interpret the version history of a Delta table



Data Preparation and Feature Engineering



Data Preparation for Machine Learning



Learning objectives

Things you'll be able to do after completing this module

- Explain data preparation and data splitting for ML, including train-test-holdout and cross-validation.
- Effectively handle missing values and understand the importance of indicator variables.
- Manage categorical features using encoding methods, considering efficiency and associated risks.
- Describe the benefits of feature standardization and recognize challenges in result interpretation.



Learning objectives

Things you'll be able to do after completing this module

- Build a data imputation and transformation pipeline for diverse datasets.
- Develop an advanced feature engineering pipeline, along with the process of saving it for future use.






Data Preparation and Feature
Engineering
LECTURE

Fundamentals of Data Preparation and Feature Engineering



The Machine Learning Full Lifecycle

Your Model is ready, then what?




Data Preparation for ML projects

Goal: Optimize input quality for accurate model predictions

Data preparation includes the following tasks:


- **Cleaning and formatting data:** This includes tasks such as **handling missing values** or outliers, ensuring data is in the correct format, and removing unneeded columns.
- **Feature Engineering:** This includes tasks like numerical **transformations**, **aggregating data**, encoding text or image data, and **creating new features**.



Feature Engineering

Transforming raw data into model-friendly features


- It is a critical step in preparing data for machine learning models.
- It transforms raw data into a format that allows ML models to extract meaningful patterns and make accurate predictions.



Feature Extraction

Transforming raw data into a set of features that better represent the underlying patterns


- Transforming Raw Data for Enhanced Modeling
- Dimensionality Reduction for Improved Performance
- Simplifying Feature Engineering



Feature Selection: Dropping Columns/Rows


Enhancing Model Efficiency

Streamlining the dataset by selectively removing columns (features) or rows (instances) to improve model efficiency and effectiveness.



Why Do We Need **Splitting Data**


Optimizing Model Evaluation



Splitting Data into Multiple Sets

Optimizing Model Training, Validation, and Testing

Splitting the dataset is a fundamental step in data preparation to facilitate effective model training, validation, and testing.



Sampling Methods

Sampling methods play a crucial role in selecting subsets of data for training, validation, and testing, influencing the performance and generalization of machine learning models. Always take care in how you are performing your sampling for your particular dataset.




Sampling Methods

Sampling methods play a crucial role in selecting subsets of data for training, validation, and testing, influencing the performance and generalization of machine learning models. Always take care in how you are performing your sampling for your particular dataset.

Here's an overview of common sampling methods:

Random Sampling:



This is the most common and basic sampling technique in machine learning. It randomly selects data points from a dataset without following any specific pattern, assuming that each data point has an equal probability of being chosen.

Example with a pandas DataFrame

```
# my_data is a pandas DataFrame and we want to sample 50  
data points  
  
my_random_sample = my_data.sample(n=50)
```




Sampling Methods

Sampling methods play a crucial role in selecting subsets of data for training, validation, and testing, influencing the performance and generalization of machine learning models. Always take care in how you are performing your sampling for your particular dataset.

Here's an overview of common sampling methods:

Stratified
Sampling:



This sampling method is used when a dataset consists of subgroups, ensuring that samples are taken from each. The population is divided into homogeneous subgroups, known as strata, and an appropriate number of instances is selected from each stratum to ensure the sample accurately represents the entire population.

Example with a pandas DataFrame & scikit-learn

```
# my_data is a pandas DataFrame
# group_label is the column containing subgroups

from sklearn.model_selection import train_test_split

stratified_train, stratified_test = train_test_split(my_data, test_size =0.2, stratify =
my_data['group_label'])
```




Sampling Methods

Sampling methods play a crucial role in selecting subsets of data for training, validation, and testing, influencing the performance and generalization of machine learning models. Always take care in how you are performing your sampling for your particular dataset.

Here's an overview of common sampling methods:

Cluster Sampling:



Cluster sampling is a sampling technique where the population is divided into naturally occurring groups, or clusters, and a subset of these clusters is randomly selected for study. Instead of sampling individuals directly, researchers collect data from all individuals within the chosen clusters.

Example with a pandas DataFrame and numpy

```
# my_data is a pandas DataFrame
# my_col is the column containing subgroup identifiers
# Randomly select M clusters from the range 1 to N

import pandas as pd
import numpy as np

clusters = np.random.choice(np.arange(1, N+1), size=M, replace=False)
cluster_sample = my_data[my_data['my_col'].isin(clusters)]
```




Cross-Validation

Cross-validation is a resampling technique used to assess the performance of a machine learning model. It involves partitioning the dataset into subsets to train and evaluate the model multiple times, providing a more robust estimate of its generalization performance.


Here are the key elements of cross-validation

- K-Fold Cross-Validation
- Stratified K-Fold Cross-Validation
- Shuffle Split Cross-Validation
- Nested Cross-Validation



Sampling for Time-Series Data

Sampling for **time-series data** requires special consideration due to the temporal dependencies inherent in the data. **Traditional random sampling or shuffling may not be appropriate**, as the order of events in time is crucial.





Data Preparation and Feature
Engineering
LECTURE



Data Imputation



Data Imputation

Data imputation is **the process of filling in missing values** in a dataset with estimated or predicted values.


The goal of data imputation is to enhance the quality and completeness of the dataset, ultimately improving the performance and reliability of the machine learning model.



Problems with Missing Data

Impacting the performance and reliability of ML models

- Reduced Model Performance
- Biased Inferences
- Imbalanced Representations
- Increased Complexity in Model Handling




The screenshot shows the Databricks Glance interface for a dataset named 'bam'. The 'Glimpse' tab is active, displaying a table with 891 rows and 12 columns. The table includes columns for PassengerId, Survived, Pclass, Name, Sex, Age, SibSp, Parch, Ticket, Fare, Cabin, and Embarked. The 'Unique values' and 'Missing values' columns are highlighted in blue and green respectively. The 'Missing values' column shows the percentage of missing data for each column: Age (19.9%), Cabin (77.1%), and Embarked (0.2%).

Column	Data type	Unique values	Missing values
PassengerId	int64	891	0
Survived	int64	2	0
Pclass	int64	3	0
Name	object	891	0
Sex	object	2	0
Age	float64	88	177 - 19.9%
SibSp	int64	7	0
Parch	int64	7	0
Ticket	object	681	0
Fare	float64	248	0
Cabin	object	148	687 - 77.1%
Embarked	object	4	2 - 0.2%



How to Handle Missing Data

Data imputation methods



Replacing Missing Values

Data imputation methods

Mean – Mode Imputation

Before	After
10	10.0
15	15.0
-	18.3
20	20.0
25	25.0

K-Nearest Neighbors (KNN with k=2)

Before	After
8	8.0
-	10.0
12	12.0
15	15.0
-	13.0

Multiple Imputation (Regression)

F1	F2	Before	After
X	X	10	10.0
X	X	15	15.0
X			$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon$
X	X	20	20.0
X	X	25	25.0



Factors Influencing Imputation Method

There isn't a definitive best approach

Nature of Data

- Is data type is **continuous, ordinal or categorical**?
- Data **distribution**: For example, median imp. Might work better for non-normal distributions

Amount of Missing Data

- How much of data is missing?
- Will imputing data improve the quality of dataset or will it add more bias?
- Multiple imp. and KNN might better for large missingness.

Form of Missingness

- Is data is **missing at random** or is there a **missingness pattern**.
- Domain knowledge and how would imputation affect downstream analysis.



Marking Imputed Data (*Best Practice*)

Keep track of imputed data

Important to mark imputed data, for:

- Model Evaluation
- Data Quality Assessment
- Enabling Transparency of Dataset
- Error Identification

ID	Name	Age	Age_imputed
1	Alice	25.0	0
2	Bob	30.0	0
3	Charlie	26.0	1
4	David	28.0	0
5	Eva	22.0	0





Data Preparation and Feature
Engineering
LECTURE

Data Encoding



Data Encoding


Why Encoding?

Data encoding is an important pre-processing step in **preparing categorical data for machine learning algorithms**, as vast majority of algorithms accept numerical input exclusively.

Issues with classic ML

- Handling high-cardinality features
- Introducing unintended relationships
- Overfitting
- Increased computational cost
- Possible lack of interpretability

Encoding Process



Working with Categorical Features

High Cardinality

Issue: Large set of categories

Many unique values in a categorical feature can lead to a large number of dummy variables, increasing dimensionality and potentially causing issues. This is known as the **high-cardinality problem**.

Possible Solutions:

Group Rare Categories:

ID	Category		ID	Category	Cat_Group
0	A		0	A	A
1	B		1	B	Rare
2	A		2	A	A
3	C		3	C	Rare
4	D	→	4	D	D
5	D		5	D	D
6	D		6	D	D
7	D		7	D	D

Top-N Categories:

ID	Category		ID	Category	Cat_Group
0	A		0	A	A
1	B		1	B	Other
2	A		2	A	A
3	C		3	C	Other
4	D	→	4	D	D
5	D		5	D	D
6	D		6	D	D
7	D		7	D	D



Working with Categorical Features

Missing Values

Issue: Categorical gaps

Categorical features often have **missing values**, which need to be addressed before model training.

Possible Solutions:

Imputation:

ID	Feature A	Feature B	ID	Feature A	Feature B
0	1.0	10.0	0	1.0	10.0
1	2.0	NaN	1	2.0	32.5
2	NaN	30.0	2	3.0	30.0
3	4.0	40.0	3	4.0	40.0
4	5.0	50.0	4	5.0	50.0

Consider Missing as a Separate Category:

ID	Feature A	Feature B	ID	Feature A	Feature B
0	1.0	10.0	0	1.0	10.0
1	2.0	NaN	1	2.0	-1.0
2	NaN	30.0	2	-1.0	30.0
3	4.0	40.0	3	4.0	40.0
4	5.0	50.0	4	5.0	50.0



Working with Categorical Features

Encoding Categories

Issue: String types

Models require numerical input, and categorical variables need to be **encoded**.

Possible Solutions:

One-Hot Encoding

ID	Category	ID	Cat_A	Cat_B	Cat_C
0	A	0	1	0	0
1	B	1	0	1	0
2	A	2	1	0	0
3	C	3	0	0	1
4	A	4	1	0	0

Label Encoding:

ID	Category	ID	Category	Label
0	A	0	A	0
1	B	1	B	1
2	A	2	A	0
3	C	3	C	2
4	A	4	A	0

Ordinal Encoding:

ID	Category	ID	Category	Ordinal
0	A	0	A	0.0
1	B	1	B	1.0
2	A	2	A	0.0
3	C	3	C	2.0
4	A	4	A	0.0




One-hot Encoding

Represent categorical variables as binary vectors

Procedure:

1. Add binary columns (0 or 1) for each unique value in categorical columns.
2. Each category is now a column that map to either 0 or 1 in reference to the ID column

Example:



ID	City
1	New York
2	Paris
3	London
4	Paris
5	New York



ID	New York	Paris	London
1	1	0	0
2	0	1	0
3	0	0	1
4	0	1	0
5	1	0	0




One-hot Encoding

Represent categorical variables as binary vectors

Drawbacks:

1. Induces sparsity
2. Limits split options for tree-based models
3. Inefficiency with high-cardinality variables
4. Obscures feature importance
5. Increases computational costs

Example:



ID	City
1	New York
2	Paris
3	London
4	Paris
5	New York



ID	New York	Paris	London
1	1	0	0
2	0	1	0
3	0	0	1
4	0	1	0
5	1	0	0



Label Encoding for Ordinal Features

Convert categorical data into numerical labels (*aka "String Indexing"*)

Procedure:

1. **Assign numeric labels:** Map each category to a numeric value based on its natural order.
2. **Transform the Feature:** Replace each categorical value in the feature column with its corresponding numeric label.

Example:

Freshman → 1
Sophomore → 2
Junior → 3
Senior → 4

ID	High School Grade Level	Age
1	Freshman	14
2	Senior	17
3	Junior	16
4	Freshman	15
5	Sophomore	16



ID	High School Grade Level	Age
1	1	14
2	4	17
3	3	16
4	1	15
5	2	16



Label Encoding for Ordinal Features

Convert categorical data into numerical labels (*aka “String Indexing”*)

Drawbacks:

- 1. Arbitrary Assignments Can Mislead Models.
- 2. Misinterpretation of Relationships.
- 3. Model Sensitivity.
- 4. Loss of Information.
- 5. Not Suitable for all algorithms.

Example:

Freshman	1
Sophomore	2
Junior	3
Senior	4

ID	High School Grade Level	Age
1	Freshman	14
2	Senior	17
3	Junior	16
4	Freshman	15
5	Sophomore	16



ID	High School Grade Level	Age
1	1	14
2	4	17
3	3	16
4	1	15
5	2	16



Target Encoding

Encoding categorical variables: replacement by mean of target variable

Procedure:

1. **Compute the mean** of the target variable for each category.
2. **Substitute** category instances with their mean values.

Example:

City	Clicked Values	Mean
New York	[1,0]	0.5
Paris	[0,1]	0.5
London	[1]	1

ID	City	Clicked
1	New York	1
2	Paris	0
3	London	1
4	Paris	1
5	New York	0



ID	City	Clicked
1	0.5	1
2	0.5	0
3	1.0	1
4	0.5	1
5	0.5	0



Target Encoding

Encoding categorical variables: replacement by mean of target variable

Drawbacks:

1. Overfitting risk
2. Handling new categories
3. Dependency of target distribution
4. Limited applicability to small datasets
5. Risk of bias

Example:

City	Clicked Values	Mean
New York	[1,0]	0.5
Paris	[0,1]	0.5
London	[1]	1

Can lead to overconfident prediction

ID	City	Clicked
1	New York	1
2	Paris	0
3	London	1
4	Paris	1
5	New York	0




ID	City	Clicked
1	0.5	1
2	0.5	0
3	1.0	1
4	0.5	1
5	0.5	0



What are Embeddings?

Embeddings: Representing Data in a Meaningful Way


- Embeddings are **dense vector representations** of data, capturing semantic relationships.
- Unlike traditional encoding methods, embeddings **place similar items closer together** in a lower-dimensional space.
- Used widely in **text, categorical data, images, and audio processing**.



How Embeddings Work?

The Role of Embeddings in Machine Learning

- Embeddings transform high-dimensional categorical or textual data into a **compact, dense vector space**.
- These representations capture relationships and context among different entities.
- Used in **Recommendation Systems, NLP, Image Search, and more**.
- Can be learned from data using neural networks or retrieved from **pretrained models** (e.g., Word2Vec, FastText).




Categorical Embedding

Transforming categorical variables into meaningful, continuous vectors

- Categorical embeddings replace sparse categorical encoding methods with lower-dimensional, dense vector representations.
- These embeddings capture relationships between different categories based on learned patterns.

Process:

1. Convert categorical variables into unique indices.
2. Pass through an embedding layer (neural network or learned lookup table).
3. The model assigns dense vector representations to each category based on learned relationships.



Embeddings vs. One-Hot Encoding

Choosing the Right Encoding for Categorical Data

Feature	One-Hot Encoding	Embeddings
Representation	Sparse binary vectors (high-dimensional)	Dense numerical vectors (low-dimensional)
Semantic Relationships	Does not capture relationships between categories	Places similar categories closer in vector space
Scalability	Inefficient for high-cardinality data	Scales well with large category sets
Efficiency	Inefficient for high-cardinality features (sparse matrix)	Efficient for high-cardinality features
Model Suitability	Suitable for simple models like Decision Trees	Best for Neural Networks and deep learning models
Example:	TV → [1, 0, 0, 0] Laptop → [0, 1, 0, 0] Phone → [0, 0, 1, 0]	TV → [0.6, 1.2, -0.8] Laptop → [0.5, 1.1, -0.7] Phone → [0.2, -0.4, 1.5]





Data Preparation and Feature
Engineering

LECTURE


Data Standardization



Data Standardization

Ensuring data consistency and optimal model performance.

Data standardization is the process of creating standards and transforming data taken from different sources into a consistent format that adheres to the standards.



Benefits of Data Standardization

Standardizing features offers significant advantages




Improved Interpretability



Prevention of Model Biases



Enhanced Convergence



Effective Regularization



Facilitates Comparison of Coefficients



Benefits of Data Standardization

Standardizing features offers significant advantages




Improved Interpretability



Prevention of Model Biases




Enhanced Convergence



Effective Regularization



Facilitates Comparison of Coefficients



Benefits of Data Standardization

Standardizing features offers significant advantages




Improved Interpretability



Prevention of Model Biases



Enhanced Convergence




Effective Regularization




Facilitates Comparison of Coefficients

Feature Influence Before Standardization



Feature Influence After Standardization



Benefits of Data Standardization

Standardizing features offers significant advantages




Improved Interpretability



Prevention of Model Biases




Enhanced Convergence



Effective Regularization



Facilitates Comparison of Coefficients



Benefits of Data Standardization

Standardizing features offers significant advantages




Improved Interpretability



Prevention of Model Biases




Enhanced Convergence



Effective Regularization



Facilitates Comparison of Coefficients



Benefits of Data Standardization

Standardizing features offers significant advantages




Improved Interpretability



Prevention of Model Biases




Enhanced Convergence



Effective Regularization



Facilitates Comparison of Coefficients



Which Models are Sensitive or Robust to the Scale of the Feature?

Sensitivity or robustness of ML models to the scale of features can vary

Models Sensitive to Feature Scale:

- Linear Models
- Support Vector Machines (SVM)
- K-Nearest Neighbors (KNN)
- Neural Networks

Models Robust to Feature Scale:

- Decision Trees and Random Forests
- Ensemble Models
- Tree-based Models
- Principal Component Analysis (PCA)



Process of Data Standardization

Navigating the Steps to Standardize Data for Machine Learning

1

Assess feature distributions/scales and need to standardize

2

Prep-process + Hold-Out split (train/test then train/validation)

3

Create/Push-to Feature Store table(s)

4

Create standardization pipeline and fit to train/validation & test sets



Drawbacks of Data Standardization

Limitations in the Process of Data Standardization

- Loss of Original Interpretation
- Challenges in Interpreting Interaction Terms
- Challenges in Result Interpretation
- *Assumption of Normally Distributed Features* (i.e. z-score)*






Data Preparation and Feature Engineering

DEMONSTRATION

Data Imputation and Transformation Pipeline



Demo

Outline

What we'll cover:

- Coerce columns to the correct data types based on feature and target requirements
- Identify and remove invalid or erroneous values in numeric columns
- Analyze missing data and drop columns or rows based on defined thresholds
- Impute boolean and string missing values using appropriate default strategies
- Understand how Spark ML handles categorical missing values using StringIndexer
- Encode categorical features using a StringIndexer → OneHotEncoder → VectorAssembler pipeline
- Apply ordered indexing for ordinal categorical features
- Split the dataset into training and test sets for modeling
- Standardize numeric features using a scaler fit on the training set only
- Impute numeric missing values using a strategy fit on the training set to avoid data leakage






Data Preparation and Feature Engineering

DEMONSTRATION

Build a Feature Engineering Pipeline with Embeddings



Demo

Outline

What we'll cover:

- Build a structured feature engineering pipeline that includes multiple preprocessing steps.
- Encode categorical features using **StringIndexer** and **OneHotEncoder**.
- Impute missing numerical values using ``Imputer`` and scale features with ``StandardScaler``.
- Assemble transformed numerical and encoded categorical features into a single feature vector.
- Prepare the target column for machine learning by converting it to a numeric label.
- Apply the feature engineering pipeline to both training and test datasets.
- Save a data preparation and feature engineering pipeline to Unity Catalog for potential future use.





Data Preparation and Feature Engineering

LAB EXERCISE

Build a Feature Engineering Pipeline



Lab

Outline


What you'll do:

- Load Dataset and Perform Data Preparation
 - Load Dataset
 - Data Preparation – Type Casting, Missing Values, Outlier Handling, Save Silver Table
- Split Dataset into Training and Test Sets
- Build Feature Engineering Pipeline
 - Analyze Data Types and Missing Values
 - Define Pipeline Stages and Construct Pipeline
- Fit the Pipeline on Training Data
- Transform Datasets and Prepare Target Column
- Save and Load the Pipeline for Reuse





Feature Store



Data Preparation for Machine Learning



© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Learning objectives

Things you'll be able to do after completing this module

- Understand Feature Store concepts and its benefits/importance in machine learning (*including improved collaboration and data consistency*).
- Explain the relationship between Feature Store tables and Unity Catalog (*and compare legacy/workspace feature store and Unity Catalog feature store*)
- Using Feature Store for Feature Engineering by creating a feature store table





Feature Store

LECTURE

Introduction to Feature Store



What is a Feature Store?

A **feature store** manages features, or input data to a machine learning model.

In a model that predicts **customer churn**, for example, features could be:

- Aggregations of raw data over time windows, like **trailing 7-day purchases**
- Joined **combinations of data sets**, like customer demographic information joined to transaction features
- Complex functions of customer information, like **estimated customer lifetime value**


The process of creating these values from data is **feature engineering**.



What is a Feature Store?

An example of a recommendation system.

A **feature store** manages features, or input data to a machine learning model.



The process of creating these values from data is **feature engineering**.



Why Would You Need a Feature Store?

Basic Motivations

Discovery



Multiple Data Scientists are trying to solve similar modeling tasks and come up with different definitions of the same features. **How can I find the features?**

Lineage



Model governance requires documentation of the features used to train a model, as well as the **upstream lineage** of a feature to reliably use it. **How is it computed, and who owns it?**

Skew



When multiple teams manage feature computation and ML models in production, minor yet significant **skew in upstream data** at the input of a feature pipeline can be very hard to detect and fix.

Online Serving




During exploration and model experimentation phases features are implemented in frameworks that do not scale to production.



Databricks for Machine Learning

A built-in and collaborative solution for the full ML lifecycle



Databricks Feature Store

Featurization

Define reusable, shareable featurization logic

Feature 1

Feature 2



save



Customer Features



Features



Feature Tables

Delta Lake based: SQL, ACLs, versions, and performance optimizations

Training Data Set Creation




Batch Scoring



Online Serving

(Lakebase Synced Tables)



Databricks Model Serving

snapshot

load

sync



Complete Integration–FS with **Unity Catalog**

Any table can be a feature table

- Feature Tables become regular UC Tables with additional metadata.
- Shared properties are unified.
 - Feature table description == table comment.
 - Feature table schema == table schema
- Three-level namespace convention

The screenshot displays the Databricks Unity Catalog interface. On the left, the 'Catalog' sidebar shows a hierarchical view of the namespace: 'labuser13154977_1774462...' (30.5 GB, 4 Cores) > 'dbacademy' > 'information_schema' > 'labuser13154977_1774462029' > 'telco_customer_features' (selected). The main panel shows the 'telco_customer_features' table details. The 'Overview' tab is active, displaying the table's description: 'Telco customer features'. Below the description is a 'Filter columns...' search bar. A table lists the columns:

Column	Type	Comment	Tags
customerID	string	PK	
gender	string		
SeniorCitizen	int		
Partner	string		
Dependents	string		



Unified Permission Model

Secure features with built-in governance

- Feature data and metadata are governed by the Unity Catalog permission model.
- Future improvements to data governance in Unity Catalog will apply to Feature data.

Catalog Explorer > dbacademy > labuser13154977_1774462029 >

telco_missing_bronze

Use with BI tools

Share

Create

Overview

Sample Data

Details

Permissions

Policies

History

Lineage

Insights

Quality

Grant

Revoke

Privileges

Inherited


Type to filter by principal

Principal	Privilege	Object
labuser13154977_1774462029@vocareum.com	ALL PRIVILEGES	dbacademy.labuser13154977_1774462029



Unified Data Lineage

Feature Store Lineage and Unity Catalog lineage become one



Integration with MLflow and Model Serving

Makes model deployment easier

Integration with MLflow

When a features from FS are used for training, **the model is packaged with feature metadata.**

In inference time, the model can look up features at runtime.

Integration with Model Serving

When the model is used for batch scoring or online inference, **it automatically retrieves features from Feature Store.**






Feature Store

DEMONSTRATION

Using Feature Store for Feature Engineering



Demo

Outline

What we'll cover:

- Create a Feature Store table from a PySpark DataFrame for training/modeling data and holdout data
- Identify the requirements for a Delta table in Unity Catalog to be automatically configured as a feature table
- Alter an existing Delta table in Unity Catalog so that it can be used as a feature table
- Add new features to an existing feature table in Unity Catalog
- Explore a Feature Store table in the user interface
- Upgrade a workspace feature table to a Unity Catalog feature table





Feature Store

LAB EXERCISE

Feature Engineering with Feature Store



Lab

Outline

What you'll do:

- Load and Prepare Dataset for Feature Selection
- Create a Feature Table
- Explore Feature Table through UI
- Access Feature Table Information
- Create Feature Table from Existing UC Table
- Enhance Feature Table with New Features
- Delete a Feature Table



Course Summary and Next Steps



Data Preparation for Machine Learning



