




# Machine Learning Model Development

---

**Databricks Academy**



© Databricks 2026. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

# Course Learning Objectives

- Describe fundamental concepts of developing ML models
- Describe the main components of MLflow for model development
- Describe hyperparameter tuning and leverage Optunta for HPO
- Utilize MLflow for model tracking and development
- Leverage the AutoML API and Databricks UI for AutoML to run an experiment, identify the model based on a preferred metric, and modify a generated model



# Learning Goals

Upon completion of this content, you should be able to:

- Describe fundamental concepts of machine learning.
- Describe the main components of MLflow for model development.
- Describe hyperparameter tuning and methods.
- Utilize MLflow for model tracking and model tuning with Optuna.
- Create an AutoML experiment, identify the best model, and modify the generated models.



# Prerequisites/Technical Considerations

Things to keep in mind before you work through this course

## Prerequisites

- 1 Knowledge of fundamental concepts of regression, classification, and clustering methods.
- 2 Familiarity with Databricks workspace and notebooks.
- 3 Intermediate level knowledge of Python.

## Technical Considerations

- 1 A cluster running on **DBR ML 17.3**
- 2 **Unity Catalog** enabled workspace



# AGENDA

| 1. Model Development Workflow           | DEMO | LAB |
|-----------------------------------------|------|-----|
| Model Development and MLflow            | ✓    | ✓   |
| Evaluating Model Performance            |      |     |
| 02. Hyperparameter Tuning               |      |     |
| Hyperparameter Tuning Fundamentals      |      |     |
| Hyperparameter Tuning with Optuna       | ✓    | ✓   |
| 03. AutoML                              |      |     |
| Automated Model Development with AutoML | ✓    | ✓   |





# Machine Learning Model Development

---

**Machine Learning Model Development**



# Problem Statement

## Things you'll be able to do after completing this module

- Describe principle categories of machine learning (supervised, unsupervised, reinforcement learning, etc.).
- Describe principle types of supervised learning: regression, classification, forecasting.
- Describe traditional steps of the machine learning model training workflow.
- Describe the benefits of MLflow on Databricks
- Describe Databricks Autologging.
- Review of key metrics for ML algorithms (regression, classification, clustering)





Model Development Workflow

LECTURE

# Model Development and MLflow






# Types of Machine Learning

## The basics of supervised, unsupervised, and reinforcement learning

### Supervised Learning

- Data has a set of input records with **associated labels**.
- Goal is to predict the output labels given an **unlabeled input**.
- Model the conditional probability:  
 $P(Y | X_1, X_2, \dots, X_n)$



### Unsupervised Learning

- Unlabeled data (no known function output)
- Clustering (group records with similar features)
- Dimensionality reduction (reduce feature space)
- Model the joint probability:  $P(X_1, X_2, \dots, X_n)$




# Types of Machine Learning

## The basics of supervised, unsupervised, and reinforcement learning


### Hybrids

- Combines supervised learning and unsupervised learning
- Labeled and unlabeled data, mostly unlabeled
- Semi-supervised:
  - Use small set of labeled data to infer labels to a larger set
- Self-supervised:
  - Learn patterns of data in one (larger) dataset that can be applied to another



### Reinforcement Learning


- States, actions, and rewards
- Useful for exploring spaces and exploiting information to maximize expected cumulative rewards
- Frequently **utilizes neural networks and deep learning**



# Types of Supervised Learning

## Regression

- Predicting a **continuous** output
- **Example:** predict sales, predict number of viewers



## Classification

- Predicting a **categorical/discrete** output
- **Example:** yes/no, image classification, predict disease




## Forecasting

- Predicting future values based on **historical** data
- **Example:** demand forecasting




# The Machine Learning Lifecycle

Model development and evaluation



# Core Machine Learning Dev. **Issues**

Modern ML lifecycle comes with many challenges



## **Reproducibility and Experiment Tracking**


Difficulty in ensuring  
consistent results across  
different environments and  
effectively documenting and  
organizing experiments






# Core Machine Learning Dev. **Issues**

Modern ML lifecycle comes with many challenges



## Reproducibility and Experiment Tracking

Difficulty in ensuring consistent results across different environments and effectively documenting and organizing experiments.




## Model Evaluation and Comparison

Challenges in accurately comparing and selecting the best-performing model due to variations in evaluation metrics and design.




# Core Machine Learning Dev. **Issues**

Modern ML lifecycle comes with many challenges



## Reproducibility and Experiment Tracking

Difficulty in ensuring consistent results across different environments and effectively documenting and organizing experiments.



## Model Evaluation and Comparison

Challenges in accurately comparing and selecting the best-performing model due to variations in evaluation metrics and design.



## Deployment and Standardization


Complexities in transitioning model, ensuring seamless integration, standardized deployment processes.



# MLflow

MLflow addresses common ML development issues.


- Open-source platform for machine learning lifecycle
- Operationalizing machine learning
- Developed by Databricks
- Pre-installed on the Databricks Runtime for ML





# MLflow Components

The four components of MLflow




APIs: CLI, Python, R, Java, REST



# MLflow Model Tracking

Make your ML workflow more manageable and transparent

- Record experiment **parameters** such as hyperparameters and model configurations.
- Log **metrics** and compare them for performance measures.
- Store and manage output **artifacts** such as trained models and plots.
- Store model's **source code** from the run.



# Model Tracking and Custom Logging

Supports for custom logging

Custom logging for;

- **Parameters** (such as custom model hyperparameters),
- Custom **metrics**
- **Artifacts** (such as the trained model and additional files) within the context of the run.

```
1 # Log metrics
2 y_pred = model.predict(X_test)
3 super_score = custom_score(y_test, y_pred)
4 mlflow.log_metric('super_score', super_score)
5
6 # Log artifacts (e.g., model file)
7 model_path = 'random_forest_model.pkl'
8 mlflow.sklearn.log_model(model, model_path)
9 mlflow.log_artifact(model_path)
10
11 # Log params
12 mlflow.log_params({'param1': 123, 'param2': 'abc'})
```



# Model Tracking and Autologging

Ensure reproducibility

Reproduce Run

Model registered

## Model Tracking Demo [Provide feedback](#)

Overview Model metrics System metrics Artifacts

Inspect, Visualize and Compare Metrics

Artifacts (Figures, datasets, text files)

```
mlflow.autolog()
```

Automatically track ML development with one line of code: parameters, metrics, data lineage, model, and environment.

Model Details

|                   |                                                                      |
|-------------------|----------------------------------------------------------------------|
| Run ID            | 70e02bb8bec2418197fc6033c0484342                                     |
| Duration          | 13.2s                                                                |
| Datasets used     | <a href="#">dataset (98736df7)</a> Training +2                       |
| Tags              | spa... : path=dbfs:/mnt/dbacademy-datasets/m... <a href="#">edit</a> |
| Source            | <a href="#">1.2 - Model Tracking with MLflow</a>                     |
| Logged models     | <a href="#">sklearn</a>                                              |
| Registered models | <a href="#">menaf_gul_8vg5_da.default.diabetes-predictions</a> v1    |

### Parameters (4)

| Q Search parameters |       |
|---------------------|-------|
| Parameter           | Value |
| criterion           | gini  |
| max_depth           | 50    |

### Metrics (8)

| Q Search metrics            |                  |
|-----------------------------|------------------|
| Metric                      | Value            |
| <a href="#">train_f1</a>    | 0.81349218257... |
| <a href="#">test_recall</a> | 0.6965090090...  |




# MLflow Concepts

Experiment → Runs → Metadata / Artifacts

## MLflow experiment:

- A higher-level organizational unit that **encompasses a set of runs**.
- Groups and organizes related runs, typically conducted to explore different configurations, parameters, or algorithms.
- Two options
  - Notebook Experiment: tied to the notebook
  - Workspace experiment: Manually specify the path. Multiple notebooks can create runs in it.






# MLflow Concepts

Experiment → **Runs** → Metadata / Artifacts

## MLflow run:

- Executions of data science code.
- Belongs to only one experiment.
- Contains
  - Metadata
  - Artifacts
- Has a default Py Func flavor which is wrapper around the native model.






# Examining Past Runs

## Via MLflow UI

- List and filter runs by metrics
- Built in to Databricks platform



| Table                    |                          |                      |              |          |                 | Metrics      | Tags                |
|--------------------------|--------------------------|----------------------|--------------|----------|-----------------|--------------|---------------------|
|                          |                          | Run Name             | Created      | Duration | Models          | val_f1_score | model_type          |
| <input type="checkbox"/> | <input type="checkbox"/> | unruly-ram-42        | 3 months ago | 30.7s    | Churn_Pred.../1 | 0.3576388... | random_forest       |
| <input type="checkbox"/> | <input type="checkbox"/> | silent-calf-856      | 3 months ago | 21.8s    | sklearn         | 0.3368421... | lightgbm            |
| <input type="checkbox"/> | <input type="checkbox"/> | loud-newt-435        | 3 months ago | 21.9s    | sklearn         | 0.2986767... | lightgbm            |
| <input type="checkbox"/> | <input type="checkbox"/> | luxuriant-shad-925   | 3 months ago | 21.3s    | sklearn         | 0.2926829... | lightgbm            |
| <input type="checkbox"/> | <input type="checkbox"/> | rebellious-goose-885 | 3 months ago | 21.3s    | sklearn         | 0.2495049... | decision_tree       |
| <input type="checkbox"/> | <input type="checkbox"/> | kindly-cat-503       | 3 months ago | 20.6s    | sklearn         | 0.2433460... | logistic_regression |
| <input type="checkbox"/> | <input type="checkbox"/> | adorable-kit-32      | 3 months ago | 22.0s    | sklearn         | 0.1957116... | lightgbm            |

## Via MLflow API

- MLflowClient Object
- List experiments
- Search runs by metrics, model etc.

```
1 from mlflow import MlflowClient
2
3 runs = MlflowClient().search_runs(
4     experiment_ids="0",
5     filter_string="params.model = 'random_forest'",
6     max_results=5,
7     order_by=["metrics.accuracy DESC"],
8 )
```







Model Development Workflow


LECTURE

# Evaluating Model Performance




# Purpose of Computing Evaluation Metrics

- Evaluation metrics provide a numerical representation of how well a model is doing.
- Metrics enable the comparison of different models or different versions of the same model.
- Evaluation metrics guide the fine-tuning of models.
- Metrics prove an objective basis for choosing between different models or approaches.



# How to Build and Evaluate Models?



| Dataset Type | Purpose                   | Usage                                          | Model Impact                                |
|--------------|---------------------------|------------------------------------------------|---------------------------------------------|
| Training     | Training data for model   | Minimize loss during optimization              | Direct impact on internal model parameters  |
| Validation   | Hyperparameter tuning     | Used for preventing overfitting                | Indirect impact on model parameters         |
| Test         | Evaluate generalizability | Assesses real-world performance on unseen data | No influence on model training or selection |



# Review of Common Metrics

## Supervised Learning – Regression

| Metric | Description | Values | Code |
|--------|-------------|--------|------|
|        |             |        |      |
|        |             |        |      |
|        |             |        |      |
|        |             |        |      |

[More Details: Appendix A](#)



# Review of Common Metrics

## Supervised Learning – Regression

| Metric | Description                                                                | Values                                    | Code                                                                              |
|--------|----------------------------------------------------------------------------|-------------------------------------------|-----------------------------------------------------------------------------------|
| $R^2$  | Measures how well the model explains the variation in the target variable. | Range: $(-\infty, 1]$<br>Larger is better | <pre>from sklearn.metrics import r2_score<br/>r2 = r2_score(y_true, y_pred)</pre> |
|        |                                                                            |                                           |                                                                                   |
|        |                                                                            |                                           |                                                                                   |
|        |                                                                            |                                           |                                                                                   |

[More Details: Appendix A](#)





# Review of Common Metrics

## Supervised Learning – Regression

| Metric              | Description                                                                   | Values                                    | Code                                                                                                 |
|---------------------|-------------------------------------------------------------------------------|-------------------------------------------|------------------------------------------------------------------------------------------------------|
| $R^2$               | Measures how well the model explains the variation in the target variable.    | Range: $(-\infty, 1]$<br>Larger is better | <pre>from sklearn.metrics import r2_score r2 = r2_score(y_true, y_pred)</pre>                        |
| Mean Absolute Error | Measures the average absolute difference between actual and predicted values. | Range: $[0, \infty)$<br>Lower is better   | <pre>from sklearn.metrics import mean_absolute_error mae = mean_absolute_error(y_true, y_pred)</pre> |
|                     |                                                                               |                                           |                                                                                                      |
|                     |                                                                               |                                           |                                                                                                      |

[More Details: Appendix A](#)





# Review of Common Metrics

## Supervised Learning – Regression

| Metric              | Description                                                                                                                    | Values                                    | Code                                                                                                         |
|---------------------|--------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------|--------------------------------------------------------------------------------------------------------------|
| $R^2$               | Measures how well the model explains the variation in the target variable.                                                     | Range: $(-\infty, 1]$<br>Larger is better | <pre>from sklearn.metrics import r2_score<br/>r2 = r2_score(y_true, y_pred)</pre>                            |
| Mean Absolute Error | Measures the average absolute difference between actual and predicted values.                                                  | Range: $[0, \infty)$<br>Lower is better   | <pre>from sklearn.metrics import<br/>mean_absolute_error<br/>mae = mean_absolute_error(y_true, y_pred)</pre> |
| Mean Squared Error  | Measures the average of squared differences between actual and predicted values. Squaring gives higher weight to large errors. | Range: $[0, \infty)$<br>Lower is better   | <pre>from sklearn.metrics import mean_squared_error<br/>mse = mean_squared_error(y_true, y_pred)</pre>       |
|                     |                                                                                                                                |                                           |                                                                                                              |

[More Details: Appendix A](#)



# Review of Common Metrics

## Supervised Learning – Regression

| Metric                  | Description                                                                                                                    | Values                                    | Code                                                                                                              |
|-------------------------|--------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------|-------------------------------------------------------------------------------------------------------------------|
| R <sup>2</sup>          | Measures how well the model explains the variation in the target variable.                                                     | Range: $(-\infty, 1]$<br>Larger is better | <pre>from sklearn.metrics import r2_score r2 = r2_score(y_true, y_pred)</pre>                                     |
| Mean Absolute Error     | Measures the average absolute difference between actual and predicted values.                                                  | Range: $[0, \infty)$<br>Lower is better   | <pre>from sklearn.metrics import mean_absolute_error mae = mean_absolute_error(y_true, y_pred)</pre>              |
| Mean Squared Error      | Measures the average of squared differences between actual and predicted values. Squaring gives higher weight to large errors. | Range: $[0, \infty)$<br>Lower is better   | <pre>from sklearn.metrics import mean_squared_error mse = mean_squared_error(y_true, y_pred)</pre>                |
| Root Mean Squared Error | Similar to MSE but takes the square root, making it more interpretable to the original units.                                  | Range: $[0, \infty)$<br>Lower is better   | <pre>from sklearn.metrics import mean_squared_error mse = mean_squared_error(y_true, y_pred, squared=False)</pre> |

[More Details: Appendix A](#)



# Review of Common Metrics

## Supervised Learning – Classification

| Metric               | Description |
|----------------------|-------------|
| Accuracy             |             |
| Precision            |             |
| Recall (Sensitivity) |             |
| F1 Score             |             |

|        |          | Prediction          |                     |
|--------|----------|---------------------|---------------------|
|        |          | Positive            | Negative            |
| Actual | Positive | True Positive (TP)  | False Negative (FN) |
|        | Negative | False Positive (FP) | True Negative (TN)  |


Confusion Matrix

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_true, y_pred)
```

# Review of Common Metrics

## Supervised Learning – Classification

| Metric               | Description                                                                                             |
|----------------------|---------------------------------------------------------------------------------------------------------|
| Accuracy             | Measures the proportion of correctly classified instances. It can be misleading for imbalance datasets. |
| Precision            |                                                                                                         |
| Recall (Sensitivity) |                                                                                                         |
| F1 Score             |                                                                                                         |




Confusion Matrix

```
from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_true, y_pred)
```

# Review of Common Metrics

## Supervised Learning – Classification

| Metric               | Description                                                                                             |
|----------------------|---------------------------------------------------------------------------------------------------------|
| Accuracy             | Measures the proportion of correctly classified instances. It can be misleading for imbalance datasets. |
| Precision            | Measures how many predicted positives are actually correct. Important when false positives are costly.  |
| Recall (Sensitivity) |                                                                                                         |
| F1 Score             |                                                                                                         |



Confusion Matrix


```
from sklearn.metrics import precision_score
precision = precision_score(y_true, y_pred)
```



# Review of Common Metrics

## Supervised Learning – Classification

| Metric               | Description                                                                                              |
|----------------------|----------------------------------------------------------------------------------------------------------|
| Accuracy             | Measures the proportion of correctly classified instances. It can be misleading for imbalance datasets.  |
| Precision            | Measures how many predicted positives are actually correct. Important when false positives are costly.   |
| Recall (Sensitivity) | Measures how many actual positives were correctly identified. Important when false negatives are costly. |
| F1 Score             |                                                                                                          |




Confusion Matrix

```
from sklearn.metrics import recall_score
recall = recall_score(y_true, y_pred)
```

# Review of Common Metrics

## Supervised Learning – Classification

| Metric               | Description                                                                                              |
|----------------------|----------------------------------------------------------------------------------------------------------|
| Accuracy             | Measures the proportion of correctly classified instances. It can be misleading for imbalance datasets.  |
| Precision            | Measures how many predicted positives are actually correct. Important when false positives are costly.   |
| Recall (Sensitivity) | Measures how many actual positives were correctly identified. Important when false negatives are costly. |
| F1 Score             | Measures the proportion of correctly classified instances. It can be misleading for imbalance datasets.  |



Confusion Matrix

```
from sklearn.metrics import f1_score
f1 = f1_score(y_true, y_pred)
```



# Precision, Recall, F1-Score

Which one to choose?

## Precision

Suitable when minimizing false positives is crucial, even at the expense of missing some positive instances.

**Example:** In a **spam email classifier**, high precision ensures that emails marked as spam are indeed spam, reducing the chances of incorrectly classifying important emails as spam.

## Recall

Appropriate when capturing all positive instances is vital, even if it means tolerating some false positives.

**Example:** In a **medical diagnosis system**, high recall ensures identifying all patients with a certain disease, even if it leads to a few false alarms.

## F1 Score

Useful when seeking a balance between precision and recall.

**Example:** In a search engine **ranking algorithm**, F1 score might be used to evaluate how well the algorithm retrieves relevant documents while minimizing irrelevant ones.




# Review of Common Metrics

## Supervised Learning – Classification

| Metric                           | Description                                                                                                                                                                      |
|----------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Log loss<br>(Cross-Entropy Loss) | Measures by comparing predicted probabilities to the actual class labels. Penalizes confident but incorrect predictions heavily. Lower values indicate better calibrated models. |
| ROC/AUC                          |                                                                                                                                                                                  |

```
from sklearn.metrics import log_loss  
ll_value = log_loss(y_true, y_pred)
```




# Review of Common Metrics

## Supervised Learning – Classification

| Metric                           | Description                                                                                                                                                                      |
|----------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Log loss<br>(Cross-Entropy Loss) | Measures by comparing predicted probabilities to the actual class labels. Penalizes confident but incorrect predictions heavily. Lower values indicate better calibrated models. |
| ROC/AUC                          | Measures the model's ability to distinguish between classes across different thresholds. AUC (Area Under the Curve) closer to 1 indicates a better model.                        |

```
from sklearn.metrics import roc_auc_score
auc_score = roc_auc_score(y_true, y_pred)
```




# Review of Common Metrics

## Unsupervised Learning – Clustering

| Metric           | Description                                                                       |
|------------------|-----------------------------------------------------------------------------------|
| Silhouette Score | Measures how well-defined clusters are. Higher values indicate better separation. |
| Elbow Method     |                                                                                   |


```
from sklearn.metrics import silhouette_score
From sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=2, random_state=42, n_init=10)
Labels = kmeans.fit_predict(X)
s_score = silhouette_score(X, labels)
```



# Review of Common Metrics

## Unsupervised Learning – Clustering

| Metric           | Description                                                                                                                             |
|------------------|-----------------------------------------------------------------------------------------------------------------------------------------|
| Silhouette Score | Measures how well-defined clusters are. Higher values indicate better separation.                                                       |
| Elbow Method     | Plots variance within clusters vs the number of clusters to find the elbow point where adding more clusters doesn't improve separation. |





# Review of Common Metrics

## Unsupervised Learning – Clustering

| Metric           | Description                                                                                                                             |
|------------------|-----------------------------------------------------------------------------------------------------------------------------------------|
| Silhouette Score | Measures how well-defined clusters are. Higher values indicate better separation.                                                       |
| Elbow Method     | Plots variance within clusters vs the number of clusters to find the elbow point where adding more clusters doesn't improve separation. |

1

```
from sklearn.metrics import silhouette_score
from sklearn.cluster import KMeans


k_values = range(1, 7)
```

2

```
inertia_values = [KMeans(n_clusters=k,
random_state=42, n_init=10).fit(X).inertia_ for k
in k_values]

knee_locator = KneeLocator(k_values,
inertia_values, curve="convex",
direction="decreasing")

optimal_k = knee_locator.elbow
```






Model Development Workflow

DEMONSTRATION

# Supervised Learning








Model Development Workflow

DEMONSTRATION

# Unsupervised Learning






Model Development Workflow

DEMONSTRATION

# Model Tracking with MLflow





Model Development Workflow

LAB EXERCISE

# Model Development Tracking with MLflow




# Appendix A



# Review of Some Common Metrics

## Supervised Learning – Regression

| Metric                  | Description |
|-------------------------|-------------|
| $R^2$                   |             |
| Mean Absolute Error     |             |
| Mean Squared Error      |             |
| Root Mean Squared Error |             |






# Review of Some Common Metrics

## Supervised Learning – Regression


| Metric                  | Description                                                                                      |
|-------------------------|--------------------------------------------------------------------------------------------------|
| $R^2$                   | Measures the proportion of variance explained by the model. Higher values indicate a better fit. |
| Mean Absolute Error     |                                                                                                  |
| Mean Squared Error      |                                                                                                  |
| Root Mean Squared Error |                                                                                                  |



# Review of Some Common Metrics

## Supervised Learning – Regression

| Metric                  | Description                                                                                      |
|-------------------------|--------------------------------------------------------------------------------------------------|
| $R^2$                   | Measures the proportion of variance explained by the model. Higher values indicate a better fit. |
| Mean Absolute Error     |                                                                                                  |
| Mean Squared Error      |                                                                                                  |
| Root Mean Squared Error |                                                                                                  |






# Review of Some Common Metrics

## Supervised Learning – Regression


| Metric                  | Description                                                                                      |
|-------------------------|--------------------------------------------------------------------------------------------------|
| $R^2$                   | Measures the proportion of variance explained by the model. Higher values indicate a better fit. |
| Mean Absolute Error     |                                                                                                  |
| Mean Squared Error      |                                                                                                  |
| Root Mean Squared Error |                                                                                                  |



# Review of Some Common Metrics

## Supervised Learning – Regression


| Metric                  | Description                                                                                      |
|-------------------------|--------------------------------------------------------------------------------------------------|
| $R^2$                   | Measures the proportion of variance explained by the model. Higher values indicate a better fit. |
| Mean Absolute Error     |                                                                                                  |
| Mean Squared Error      |                                                                                                  |
| Root Mean Squared Error |                                                                                                  |



# Review of Some Common Metrics

## Supervised Learning – Regression


| Metric                  | Description                                                                                      |
|-------------------------|--------------------------------------------------------------------------------------------------|
| $R^2$                   | Measures the proportion of variance explained by the model. Higher values indicate a better fit. |
| Mean Absolute Error     |                                                                                                  |
| Mean Squared Error      |                                                                                                  |
| Root Mean Squared Error |                                                                                                  |



# Review of Some Common Metrics

## Supervised Learning – Regression

| Metric                  | Description                                                                                                      |
|-------------------------|------------------------------------------------------------------------------------------------------------------|
| R <sup>2</sup>          | Measures the proportion of variance explained by the model. Higher values indicate a better fit.                 |
| Mean Absolute Error     | Computes the average of the absolute difference between actual and predicted values. Less sensitive to outliers. |
| Mean Squared Error      |                                                                                                                  |
| Root Mean Squared Error |                                                                                                                  |




$$MAE = \frac{1}{n} \sum_{i=1}^n \overset{\text{Residual}}{|y_i - \hat{y}_i|}$$



# Review of Some Common Metrics

## Supervised Learning – Regression

| Metric                  | Description                                                                                                      |
|-------------------------|------------------------------------------------------------------------------------------------------------------|
| R <sup>2</sup>          | Measures the proportion of variance explained by the model. Higher values indicate a better fit.                 |
| Mean Absolute Error     | Computes the average of the absolute difference between actual and predicted values. Less sensitive to outliers. |
| Mean Squared Error      | Penalizes larger errors more than MAE due to squaring. Common for Optimization.                                  |
| Root Mean Squared Error |                                                                                                                  |



$$MSE = \frac{1}{n} \sum_{i=1}^n (\text{Residual } y_i - \hat{y}_i)^2$$






# Review of Some Common Metrics

## Supervised Learning – Regression

| Metric                  | Description                                                                                                      |
|-------------------------|------------------------------------------------------------------------------------------------------------------|
| $R^2$                   | Measures the proportion of variance explained by the model. Higher values indicate a better fit.                 |
| Mean Absolute Error     | Computes the average of the absolute difference between actual and predicted values. Less sensitive to outliers. |
| Mean Squared Error      | Penalizes larger errors more than MAE due to squaring. Common for Optimization.                                  |
| Root Mean Squared Error | Similar to MSE but takes the square root, making it more interpretable to the original units.                    |




$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

Residual



# Bias–Variance Tradeoff

## Underfitting vs Overfitting




|          | Description of Error                                                        | Mitigation Strategies                                                                                                |
|----------|-----------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------|
| Bias     | Error due to overly <b>simplistic assumptions</b> in the model.             | Add more features or increasing model complexity.                                                                    |
| Variance | Error due to <b>sensitivity to small fluctuations</b> in the training data. | Add more data to help with generalizability. Add regularization, reduce model complexity, or using cross-validation. |





# Bias–Variance Tradeoff

## Underfitting vs Overfitting



|               | Reason                                                                 | Increase/<br>Decrease            |
|---------------|------------------------------------------------------------------------|----------------------------------|
| High Bias     | Model is too simple.                                                   | ▲ FP & ▲ FN                      |
| High Variance | Model does not generalize well, leading to precision-recall tradeoffs. | ▲ FP & ▼ FN<br>or<br>▼ FP & ▲ FN |



# Review of Some Common Metrics

## Supervised Learning – Classification

| Metric               | Description                                                                                             |
|----------------------|---------------------------------------------------------------------------------------------------------|
| Accuracy             | Measures the proportion of correctly classified instances. It can be misleading for imbalance datasets. |
| Precision            |                                                                                                         |
| Recall (Sensitivity) |                                                                                                         |
| F1 Score             |                                                                                                         |

|        |          | Prediction               |                          |
|--------|----------|--------------------------|--------------------------|
|        |          | Positive                 | Negative                 |
| Actual | Positive | True Positive (TP)<br>4  | False Negative (FN)<br>1 |
|        | Negative | False Positive (FP)<br>1 | True Negative (TN)<br>4  |

$$\frac{TP+TN}{TP+TN+FP+FN} = \frac{4}{5}$$



# Review of Some Common Metrics

## Supervised Learning – Classification

| Metric               | Description                                                                                             |
|----------------------|---------------------------------------------------------------------------------------------------------|
| Accuracy             | Measures the proportion of correctly classified instances. It can be misleading for imbalance datasets. |
| Precision            | Measures how many predicted positives are actually correct. Important when false positives are costly.  |
| Recall (Sensitivity) |                                                                                                         |
| F1 Score             |                                                                                                         |

|        |          | Prediction               |                          |
|--------|----------|--------------------------|--------------------------|
|        |          | Positive                 | Negative                 |
| Actual | Positive | True Positive (TP)<br>4  | False Negative (FN)<br>1 |
|        | Negative | False Positive (FP)<br>1 | True Negative (TN)<br>4  |

$$\frac{TP}{TP+FP} = \frac{4}{5}$$



# Review of Some Common Metrics

## Supervised Learning – Classification

| Metric               | Description                                                                                              |
|----------------------|----------------------------------------------------------------------------------------------------------|
| Accuracy             | Measures the proportion of correctly classified instances. It can be misleading for imbalance datasets.  |
| Precision            | Measures how many predicted positives are actually correct. Important when false positives are costly.   |
| Recall (Sensitivity) | Measures how many actual positives were correctly identified. Important when false negatives are costly. |
| F1 Score             |                                                                                                          |

|        |          | Prediction               |                          |
|--------|----------|--------------------------|--------------------------|
|        |          | Positive                 | Negative                 |
| Actual | Positive | True Positive (TP)<br>4  | False Negative (FN)<br>1 |
|        | Negative | False Positive (FP)<br>1 | True Negative (TN)<br>4  |

$$\frac{TP}{TP+FN} = \frac{4}{5}$$



# Review of Some Common Metrics

## Supervised Learning – Classification

| Metric               | Description                                                                                              |
|----------------------|----------------------------------------------------------------------------------------------------------|
| Accuracy             | Measures the proportion of correctly classified instances. It can be misleading for imbalance datasets.  |
| Precision            | Measures how many predicted positives are actually correct. Important when false positives are costly.   |
| Recall (Sensitivity) | Measures how many actual positives were correctly identified. Important when false negatives are costly. |
| F1 Score             | Measures the proportion of correctly classified instances. It can be misleading for imbalance datasets.  |

|        |          | Prediction               |                          |
|--------|----------|--------------------------|--------------------------|
|        |          | Positive                 | Negative                 |
| Actual | Positive | True Positive (TP)<br>4  | False Negative (FN)<br>1 |
|        | Negative | False Positive (FP)<br>1 | True Negative (TN)<br>4  |

$$2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} = \frac{4}{5}$$






# Review of Some Common Metrics

## Supervised Learning – Classification

| Metric                           | Description                                                                                            |
|----------------------------------|--------------------------------------------------------------------------------------------------------|
| Log loss<br>(Cross-Entropy Loss) | Measures how many predicted positives are actually correct. Important when false positives are costly. |
| ROC/AUC                          |                                                                                                        |


$$-\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$



# Hyperparameter Tuning

---

**Machine Learning Model Development**



# Learning objectives

Things you'll be able to do after completing this module

- Define a hyperparameters
- Describe, compare, and contrast grid, random, and Bayesian search methods for hyperparameter tuning
- Describe cross-validation as a method for robust performance estimation
- Describe Optuna as an open-source hyperparameter optimization (HPO) framework







Hyperparameter Tuning

LECTURE

# Hyperparameter Tuning Fundamentals



# What is a Hyperparameter

## Definitions

- **Hyperparameter** is a parameter whose value is used to **control** the training process.
- **Hyperparameter Tuning** (model tuning) is the process of adjusting the hyperparameters of a model **to find the best configuration** that **maximizes its performance** on a given dataset.

## Example Hyperparameters:


- Random Forest:
  - Maximum depth of trees
  - Number of trees
  - Number of features
- Linear regression
  - Normalize features or not
- Neural Networks
  - Activation function
  - Number of neurons in each hidden layer




# Hyperparameter Search Methods

How to pick a hyperparameter?


## Grid Search



## Random Search



## Bayesian Optimization




[Image Source](#)



# Grid Search

## Systematic parameter selection

- Define a grid of hyperparameter values and exhaustively searching through all possible combinations.
- Evaluate each combination using cross-validation and selects the best set of hyperparameters.
- Can be computationally expensive, especially when the number of hyperparameters and their possible values is large




Grid Search will always find the best parameter, but is computational expensive.



# Random Search

## Stochastic parameter selection

- Randomly samples hyperparameter values from specified distributions.
- Random search is more efficient than grid search, especially when dealing with large datasets or hyperparameters.
- It doesn't guarantee the best combination.



Random Search generally outperforms grid search, but does not systematically cover all possible values.






# Bayesian Optimization

## Probabilistic parameter selection

- Aiming to **minimize validation error** by constructing a **probability model** of the objective function.
- Guides the search based on **previous evaluations** to focus on promising regions of the hyperparameter space.
- **The Tree of Parzen Estimators (TPE)** algorithm is a popular Bayesian method.




Bayesian approaches generally outperform grid search and random search, but more complex to implement.



# Validation

- The validation set is one of our holdout datasets
- Build a model for each set of hyperparameter values
- Evaluate each model to identify the optimal hyperparameter values
- What dataset should we use to train and evaluate?




**What if there isn't enough data to split into three separate sets?  
Risk of overfitting due to non-random initial split?**



# Validation

- The validation set is one of our holdout datasets
- Build a model for each set of hyperparameter values
- Evaluate each model to identify the optimal hyperparameter values
- What dataset should we use to train and evaluate?




**What if there isn't enough data to split into three separate sets?  
Risk of overfitting due to non-random initial split?**





# K-Fold Cross Validation

A method for robust performance estimation



Final Pass:

Training with Optimal Hyperparameters

Test



# Tuning Can be Expensive!

## A Grid Search example

Train and validate **every unique combination** of hyperparameters

| Tree Depth | Number of Trees |
|------------|-----------------|
| 5          | 2               |
| 8          | 4               |



| Tree Depth | Number of Trees |
|------------|-----------------|
| 5          | 2               |
| 5          | 4               |
| 8          | 2               |
| 8          | 4               |

✖ K (folds)

+ Best model

**Question:** With 3-fold cross validation, how many models will this build?



# Optuna



# Popular Tools for Hyperparameter Opt.



- **Scikit-learn's** GridSearchCV and RandomizedSearchCV: Offer grid search and random search for hyperparameter tuning.



- **Hyperopt**: A popular library for performing hyperparameter optimization using Bayesian methods. **Hyperopt will be removed with DBR ML version 17.0+.**



- **Ray**: Ray Tune provides a scalable library for hyperparameter tuning, offering both grid and advanced search algorithms.



- **Optuna**: An open source hyperparameter optimization framework to automate hyperparameter search.



# Optuna: Lightweight, Versatile, and Platform Agnostic

Simple installation and minimal requirements

- Optuna is written entirely in Python with few dependencies.
- Intuitive API:
  - Define parameters within your code.
  - No need for framework-specific syntax.
- Platform and Framework Agnostic
  - e.g. PyTorch, TensorFlow, Keras, Scikit-Learn, XGBoost, LightGBM, etc.





# Optuna Terminology

## Definitions

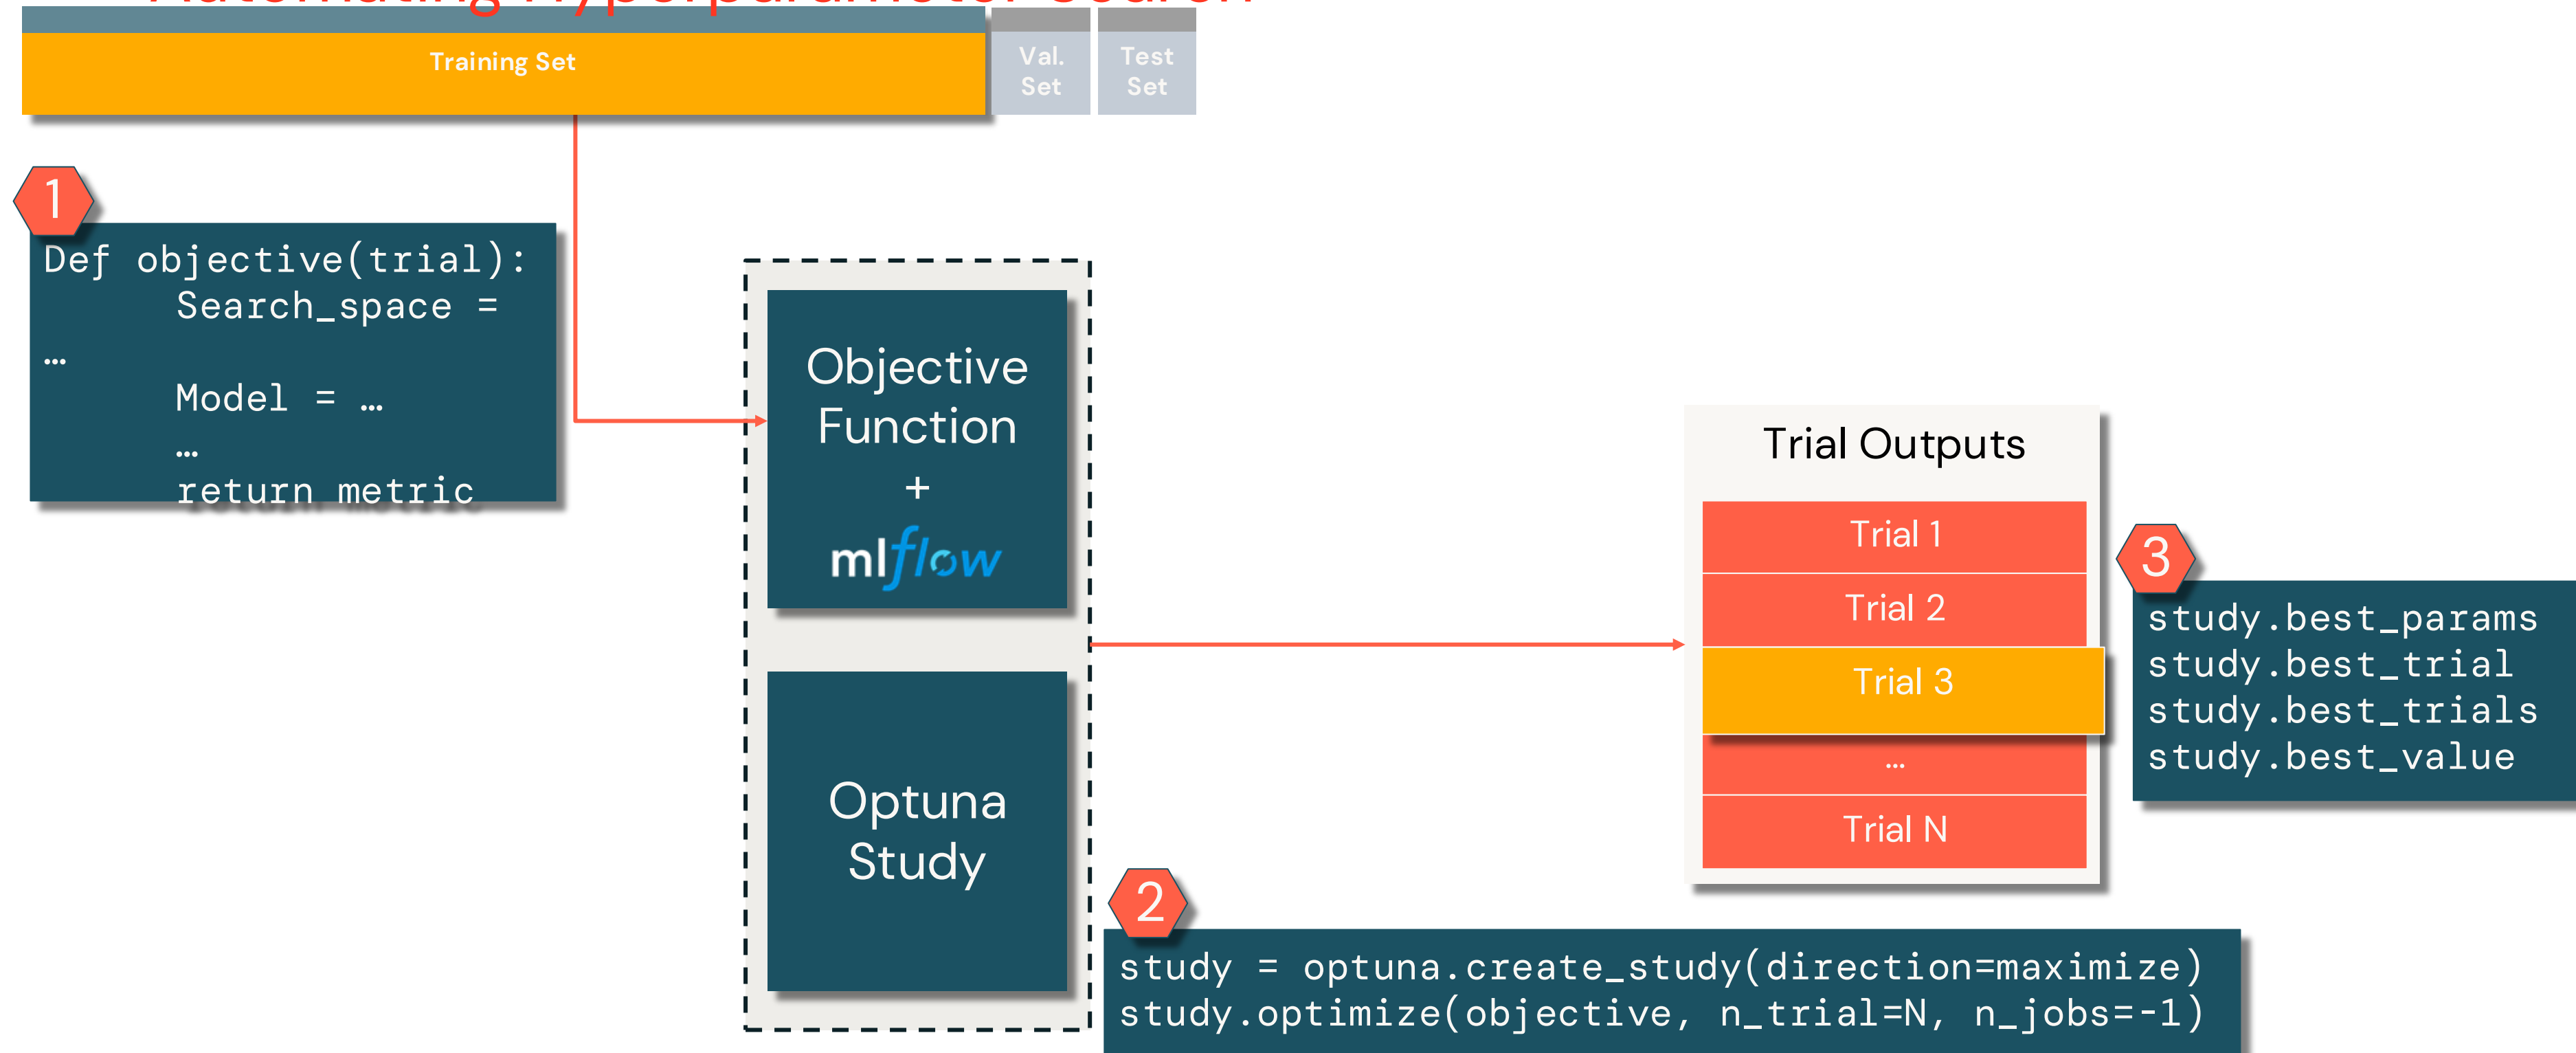
- A **trial** is a single call of the objective function.
- A **study** is an optimization session, which is a collection of trials.
- A **parameter** is a variable whose value is to be optimized, which is typically the loss or accuracy of the model.

Optuna is a framework designed for the *automation* and the *acceleration* of the *optimization* studies.



# Hyperparameter Optimization Framework

## Automating Hyperparameter Search





# Key Concepts

## Search space and algorithms

### Search Space Definition with Python Syntax

- Hyperparameter search space can be defined using functions such as;
  - `suggest_discrete_uniform`
  - `suggest_categorical`
  - `suggest_float`
  - `suggest_int`
  - `suggest_uniform`

### Sampling Strategy

- Sampling strategy determines how to decide the next hyperparameter.
- Continuously refine the search space using past evaluation results to focus on promising areas.
- Some of the supported strategies; Grid Search, Random Search, **Tree-structured Parzen Estimator algorithm (default)**

### Pruning Strategy\*

- Early stopping halts unpromising trials early based on intermediate results.
- Call `report()` and `should_prune()` after each step to activate pruning.

\*Not covered in this course.



# Optuna – MLflow Integration

Track hyperparameters and metrics of all the Optuna trials

**MLflowCallback:** This callback enables automatic logging of hyperparameters and metrics for each trial into MLflow, streamlining experiment tracking.

**Parent–Child Run Organization:** MLflow's support for hierarchical run structures allows grouping of all Optuna trials under a single parent run.

```
from optuna.integration.mlflow import MLflowCallback
mlflow_callback = MLflowCallback(
    tracking_uri = "databricks",
    create_experiment= False,
    mlflow_kwargs = {
        "experiment_id": <experiment_id>,
        "Nested": True
    }
}

study.optimize(
    objective,
    n_trials = N,
    callbacks = [mlflow_callback])
```






Hyperparameter Tuning

DEMONSTRATION

# Hyperparameter Tuning with Optuna





Hyperparameter Tuning

LAB EXERCISE

# Hyperparameter Tuning with Optuna






# AutoML

---

**Machine Learning Model Development**





# Learning objectives

Things you'll be able to do after completing this module

- Describe the benefits of using AutoML to generate machine learning models.
- Describe the glass box nature of AutoML.
- Start an AutoML experiment via the UI and the API.
- Open and edit a notebook generated by AutoML.
- Identify the best model generated by AutoML based on a given metric.
- Modify the best model generated by AutoML.





AutoML

LECTURE


# Introduction to AutoML






# What is AutoML?

**Automated machine learning (AutoML)** is a fully-automated model development solution seeking to “democratize” machine learning. While the scope of the automation varies, AutoML technologies usually **automate the ML process from data to model selection**.



# Two Key Pain Points Solved with AutoML

**Quickly Verify the Predictive Power of a Dataset**



*“Can this dataset be used to predict customer churn?”*

**Get a Baseline Model to Guide Project Direction**




*“What direction should I go in for this ML project and what benchmark should I aim to beat?”*



# Problems with Existing AutoML Solutions

## Opaque-Box and Production Cliff Problems in AutoML




| Problem                                                                                                                                                                                                                                    | Result / Pain Points                                                                                                                                                                                    |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|  Data scientists need to be able to explain how they trained a model for regulatory purposes (e.g., FDA, GDPR, etc.) - “opaque box” models do not help. |  Time and energy required to reverse engineering these “opaque-box” returned models to modify and/or explain them. |
|  A “production cliff” exists where data scientists need to modify the returned “best” model using their domain expertise.                               |  The “best” model returned is often not good enough to deploy                                                      |



# Solution: Databricks AutoML

Rapid, simplified machine learning built on the Lakehouse Architecture



# Solution: Databricks AutoML

Rapid, simplified machine learning **for everyone**

## Quick-start ML initiatives


Accelerate your time to production,  
Save weeks on ML projects

## Auto-generated notebooks

Ensure best practices.  
Customize baseline models with  
your domain expertise

## Wide range of problems

Solve classification, regression, and  
forecasting problems from a variety  
of ML libraries





# Databricks AutoML

A glass-box solution empowering data teams without taking away control

UI and API to start AutoML training

**Configure AutoML experiment** [Provide feedback](#)

1 Configure 2 Join Feature 3 Train 4 Evaluate

**Compute Configuration**

\* Cluster (Databricks Runtime 9.1 LTS ML or above)

Menaf-Dev

**Experiment Configuration**

\* ML problem type

Regression

\* Input training dataset

Browse hive\_metastore.menaf\_gul\_8vg5\_da\_gdml.customr

\* Prediction target

units\_purchased

\* Experiment name


units\_purchased\_customer\_features-2023\_05\_01-14\_43

**Advanced Configuration (optional)**

| Run Name             | Created     | Duration | Source     | Models  | Metrics |
|----------------------|-------------|----------|------------|---------|---------|
| rebellious-finch-526 | 11 days ago | 1.9min   | Noteboo... | sklearn | 0.865   |
| vaunted-calf-990     | 11 days ago | 1.9min   | Noteboo... | sklearn | 0.864   |
| smiling-lark-727     | 11 days ago | 2.2min   | Noteboo... | sklearn | 0.864   |
| Experiment 2         | 11 days ago | 1.8min   | Noteboo... | sklearn | 0.862   |
| languid-vole-290     | 11 days ago | 1.9min   | Noteboo... | sklearn | 0.862   |
| industrious-coit-828 | 11 days ago | 1.9min   | Noteboo... | sklearn | 0.862   |
| debonair-gnat-986    | 11 days ago | 1.8min   | Noteboo... | sklearn | 0.862   |
| capable-skink-256    | 11 days ago | 1.7min   | Noteboo... | sklearn | 0.862   |
| painted-bear-922     | 11 days ago | 2.0min   | Noteboo... | sklearn | 0.862   |
| fun-steed-935        | 11 days ago | 1.7min   | Noteboo... | sklearn | 0.862   |


Auto-created MLflow Experiment to track models and metrics

Easily deploy to Model Registry



Auto-generated *Data Exploration* notebook

Understand and debug data quality and preprocessing



Auto-generated notebooks with source code

Iterate further on models from AutoML, adding your expertise





# “Glass-Box” AutoML with a UI

## Configure

## Train and Evaluate

## Customize

## Deploy

The image displays four sequential screenshots of the Databricks AutoML interface, illustrating the workflow from configuration to deployment.

- Configure:** The first screenshot shows the 'Configure AutoML experiment' page. It includes a progress bar with steps: 1. Configure, 2. Augment, 3. Train, and 4. Evaluate. The 'Configure' step is active. The configuration details include:
  - Compute: dais\_mlr\_8\_new
  - ML problem type: Classification
  - Training data: default.usage\_logs
  - Target column: isMining
  - Experiment name: isMining\_usage\_logs-2021\_05-22\_33
  - Data directory: (empty)
- Train and Evaluate:** The second screenshot shows the 'AutoML Evaluation' page. It indicates that the evaluation is complete. The 'Model with best val\_f1\_score' is highlighted. A table lists 16 matching runs, showing columns for Start Time, Run Name, User, and Source. The 'Train' step is active in the progress bar.
- Customize:** The third screenshot shows the 'Generated Trial Notebook' page. It displays a Python notebook with code for Random Forest training, preprocessing, and inference. The 'Customize' step is active in the progress bar.
- Deploy:** The fourth screenshot shows the 'Model Versions' page. It displays the model's status as 'Ready - Stop' and provides the model URL. The 'Deploy' step is active in the progress bar.



# AutoML Supports Common Problem Types

Support for problem, model, feature types

## Problem Types



Classification



Regression



Time Series Forecasting  
(**Serverless**)

## Models / Tuning



**XGBoost**



HYPEROPT

 **LightGBM**

**PROPHET**

## Features



Numerical



Categorical



Timestamp



Text





# Complete Integration

AutoML integrated with other products

AutoML integrates with;

- **Feature Store:** You can expand the input training dataset using existing feature tables.
- **Model Tracking:** All trial run metrics and parameters are tracked.
- **Model Registry:** You can register generated models to Unity Catalog.
- **Model Serving:** You can deploy generated models with Model Serving for real-time inference.






AutoML

DEMONSTRATION

# Automated Model Development with AutoML



# Demo

## Outline

### What we'll cover:

- Prepare data
- AutoML Experiment with the UI
  - Create an experiment
  - View the best run
- AutoML Experiment with the API
  - Start an experiment
  - Search for best run
  - Import notebooks for other experiments





AutoML

LAB EXERCISE

# AutoML



# Lab

## Outline

### What you'll do:

- Load dataset
- Create a classification experiment using the AutoML UI
- Create a classification experiment with the AutoML API
- Retrieve the best run and show the model URI
- Import the notebook for a run





# Summary and Next Steps

---

**Machine Learning Model Development**



